




# Sentinel-2

## Land

# Classification



# **Sentinel-2 Land Classification Project Report.**

## **1. Introduction and Executive Summary:**

### **1.1. Project Overview and Goals**

This project successfully developed an automated, high-accuracy system for classifying various land cover types using advanced Deep Neural Networks (DNNs). The system leverages free, multispectral satellite images captured by the European Space Agency's **Sentinel-2** satellite. The primary objective was to categorize land into major classes, such as Agricultural Lands, Water Bodies, Urban Areas, Desert, Roads, and Trees/Forests.

The core goal was to address the limitations of traditional remote sensing methods by enabling **automatic feature extraction** from multi-band imagery, ultimately delivering a robust and production-ready land classification engine.

### **1.2. Significance and Applications**

The resulting land classification system is a vital tool designed to support a wide range of critical real-world applications:

- **Smart Urban Planning:** Providing accurate data on land use to inform development and infrastructure decisions.
- **Environmental Monitoring:** Tracking changes in ecosystems, deforestation, and water bodies over time.
- **Resource Management:** Aiding in the efficient management of natural resources, particularly agricultural lands and water bodies.
- **Disaster Response:** Quickly assessing the state of infrastructure and land post-disaster.

### **1.3. Executive Summary of Results**

The project adhered to a rigorous five-milestone structure, culminating in the deployment of a high-performance deep learning model<sup>9</sup>. The team selected the **EfficientNet-B0** architecture as the final classification engine. This choice was based on its superior efficiency and high accuracy while maintaining a low computational footprint, making it ideal for scalable deployment. The model achieved the following outstanding performance metrics:

Metric	Result
Final Test Accuracy	<b>97.12%</b>
Validation Accuracy	<b>97.00%</b>
Key AI Model	EfficientNet-B0 (Deep Neural Network)
Deployment Solution	Simple Web Interface using Flask API

## 2. Data Acquisition and Exploration

### 2.1. Data Source and Collection

The project utilized the **EuroSat Labeled Dataset**, which was derived from the Sentinel-2 mission images. This dataset was selected due to its high quality, public availability, and labeled classification across 10 distinct land types, encompassing 13 spectral bands. Initial data exploration also involved checking official platforms like the Copernicus Open Access Hub and USGS Earth Explorer, but EuroSat was chosen for its ready-made, labeled nature.

The Sentinel-2 imagery is multispectral, requiring the inclusion of bands from the Visible Spectrum (RGB) and the crucial Near-Infrared (NIR) band.

### 2.2. Land Classification Categories

The EuroSat dataset used for training, validation, and testing included images covering a broad range of land types, such as:

- Annual Crop
- Forest
- Herbaceous Vegetation
- Highway
- Industrial (Urban/Built-up)
- Pasture
- Permanent Crops
- Residential
- River (Water Bodies)
- Sea/Lake (Water Bodies)




## 2.3. Spectral Band Analysis and Feature Engineering

Exploratory Data Analysis (EDA) confirmed the critical role of specific non-visible bands (Near-Infrared and Red Edge) for differentiating between land classes<sup>21</sup>. This insight justified the decision to leverage **all 13 spectral bands** rather than relying solely on the standard RGB bands.

A key aspect of feature engineering involved calculating vegetation indices, such as the **Normalized Difference Vegetation Index (NDVI)**. This index was essential for the initial Traditional Machine Learning baseline models, which relied heavily on these manually engineered features due to their limited capacity to handle the high dimensionality of raw multispectral images.

**Heatmap of Mean Reflectance Across Classes and Bands**




### 3. Preprocessing and Feature Engineering

#### 3.1. Data Preparation Pipeline

The raw image data underwent several critical transformations to ensure optimal model compatibility and performance.

- **Standardizing Image Sizes:** All input images were resized to a consistent dimension required by the EfficientNet-B0 architecture.
- **Atmospheric Correction:** This process was applied to remove distortions caused by atmospheric scattering and absorption, providing a more accurate representation of the land surface.
- **Normalization:** Pixel values across the multispectral bands were scaled to a consistent range. This is crucial for optimal neural network training, as it prevents large input values from disproportionately affecting the learning process.
- **Data Splitting:** The collected dataset was rigorously divided into three subsets to facilitate robust model training and evaluation: Training (70%), Validation (15%), and Test (15%) sets.

EuroSAT Dataset Distribution by Land Cover Class



#### 3.2. Data Augmentation

To significantly increase the dataset diversity, enhance the model's ability to generalize, and reduce the risk of overfitting, systematic **Data Augmentation** techniques were applied exclusively to the training set.

- **Techniques Used:** This included various transformations such as **Rotation**, **Flips** (horizontal and vertical), and **Cropping**.
- **Implementation:** These transformations were managed using specialized image processing libraries like **Torch vision**, **PIL**, and **OpenCV**.


## 4. Model Selection and Experimentation

### 4.1. Traditional Machine Learning (ML) Baseline Performance

A performance benchmark was established using several traditional Machine Learning algorithms<sup>1</sup> before transitioning to Deep Learning. The bar chart (Comparison of Machine Learning Algorithm Accuracies) illustrates the performance achieved by these models on the land classification task.

- **Models Tested:** The team tested **Decision Tree**, **Random Forest**, **Support Vector Machine (SVM)**, **K-Nearest Neighbors (KNN)**, and **Logistic Regression**.
- **Input Data:** These models relied on **manually engineered features** (such as NDVI and other spectral band statistics) rather than raw pixel data, due to their limited capacity to handle the high dimensionality of multispectral images.



Model	Accuracy
Random Forest	<b>0.813</b>
SVM	0.776
KNN	0.757
Decision Tree	0.694
Logistic Regression	0.631



### 4.2. Deep Learning Architecture Rationale:

The transition to Deep Learning addressed the limitations of traditional ML by enabling **automatic feature extraction** directly from the multi-band imagery. The team tested advanced Convolutional Neural Network (CNN) architectures using a sophisticated transfer learning approach.

- **ResNet-50:** This architecture was tested for its proven ability to train very deep networks using **skip connections**, which effectively mitigates the vanishing gradient problem.
- **EfficientNet-B0 (Final Choice):** Selected as the final architecture due to its superior **performance-to-efficiency ratio**. EfficientNet leverages a **compound scaling method** to optimize model depth, width, and resolution simultaneously. It achieved the best overall performance, offering high accuracy with a relatively low computational footprint, making it ideal for scalable deployment.



## Conclusion

The **Random Forest** model achieved the highest accuracy among the traditional methods, scoring **0.813**. However, this performance was considered **suboptimal**. The main conclusion drawn was that the reliance of these models on **manual feature extraction** and their limited ability to capture complex spatial patterns necessitated the transition to more advanced techniques, specifically **Convolutional Neural Networks (CNNs)**. The final Deep Learning model, EfficientNet-B0, significantly outperformed this baseline, achieving a test accuracy of 97.12%.

## 5. Training and Optimization

### 5.1. Implementation and Strategy


The model was built and trained using the **PyTorch** deep learning framework, selected for its flexibility in custom training loops and strong community support.

- **Transfer Learning:** A pre-trained version of **EfficientNet-B0** (using ImageNet weights) was utilized, and the top layers were **fine-tuned** specifically for the EuroSat classification task.
- **Hyperparameter Tuning:** A systematic approach was used to optimize the learning rate, batch size, and the **Adam Optimizer** parameters to ensure rapid convergence and optimal performance.
- **Training Strategy:** Techniques such as **Cross-Validation** and **Early Stopping** based on the validation loss were applied during training.

### 5.2. Overfitting Prevention

To ensure the model's performance on unseen data (generalization), strategies were implemented to prevent overfitting:

- **Dropout Layers:** These were systematically included in the network.
- **Early Stopping:** This technique was implemented based on monitoring the validation loss, halting the training process when validation performance ceased to improve.



---

## ***6. Evaluation and Performance Results***

### ***6.1. Final Model Performance***

The final model demonstrated strong generalization capabilities on the Test Set, significantly outperforming the Traditional ML baseline.

Metric	EfficientNet-B0 Result
Test Accuracy	<b>97.12%</b>
Validation Accuracy	<b>97.00%</b>

### ***6.2. Evaluation Metrics***

Performance was measured and verified across multiple critical metrics to confirm robust and balanced classification across all land classes.

- **Metrics Used:** Accuracy, Precision, Recall, and F1-Score.
- **Key Insight:** The high values across all metrics confirmed that the model was not only accurate overall but also provided balanced classification without bias toward any specific class.

---

## ***7. Deployment and Monitoring (MLOps)***

### ***7.1. Model Deployment***

The finalized **EfficientNet-B0** model was prepared for a production environment to ensure its predictive power was accessible for practical use.

- **Deployment API:** The trained model was serialized and deployed as a **RESTful API** using the lightweight **Flask** web framework. This API is designed to receive a Sentinel-2 image as input and return the predicted land type.

- **User Interface (UI):** A simple, functional web interface was developed using Flask to demonstrate the model's capability, allowing users to upload an image and view the classification result immediately.
- **Scalability:** While initially deployed locally via Flask, the strategy called for future hosting on cloud platforms (Google Cloud, AWS, or Azure) to ensure scalability.

The screenshot shows a web-based application titled "Satellite AI Vision". At the top, there is a banner with the text "AI-Powered Satellite Analysis" and "Transform satellite imagery into actionable intelligence. Our advanced deep learning model analyzes land cover with exceptional accuracy, helping you understand terrain at scale." Below the banner are three buttons: "LAND CLASSES", "ACCURACY", and "INFERENCE".

The main area of the application is divided into two sections. On the left, there is a "Upload Imagery" section with a file input field labeled "Image" and a preview area showing a small dark square thumbnail. Below this is a blue button labeled "Analyze Terrain". On the right, there is a "Classification Results" section for "Forest" land cover. It displays a bar chart with the following data:

Land Cover Category	Percentage
Forest	100%
Residential	0%
Herbaceous Vegetation	0%
Permanent Crop	0%

At the bottom of the right section is a button labeled "View All Land Classes".

At the very bottom of the page, there is a call-to-action button labeled "Try Sample Images" with the text "Click any example below to instantly test the classification system".

## 7.2. MLOps Strategy

Continuous monitoring was documented as a foundational MLOps strategy to track performance and ensure the system's long-term viability.

- **Performance Tracking:** Establishing a baseline for accuracy and latency to measure future degradation.
- **Drift Detection:** A plan was established for periodically checking for **Model Drift** (decay in performance over time) and **Data Drift** (changes in incoming satellite imagery patterns).
- **Retraining Plan:** A clear strategy for periodic retraining was established to incorporate new, diverse data and maintain the target approx97% accuracy threshold over time.

## 8. Technical Stack and Libraries

The following table summarizes the core technical tools and libraries used throughout the project lifecycle:

Category	Tools and Libraries Used	Rationale
Deep Learning Framework	<b>PyTorch</b>	Chosen for its dynamic graph computation, flexibility, and strong community support for advanced research models.
Core Model Architecture	<b>EfficientNet-B0</b>	State-of-the-art model selected for its high accuracy and parameter efficiency through compound scaling.
Deployment	<b>Flask</b>	Lightweight and flexible Python web framework used to create a rapid and simple API for the model inference.
Data Handling	<b>NumPy, Pandas</b>	Essential libraries for array manipulation, data structure management, and statistical analysis of spectral data.
Image Processing	<b>Torchvision, PIL, OpenCV</b>	Used for data loading, preprocessing, augmentation, and applying required transformations.
Visualization	<b>Matplotlib, Seaborn</b>	Utilized for generating visual reports, confusion matrices, and analyzing spectral distributions.

## 9. Conclusion and Future Recommendations

### 9.1. Conclusion

The project successfully delivered a robust and efficient land classification system based on Sentinel-2 data. The comprehensive methodology, ranging from detailed preprocessing to advanced model deployment (MLOps), has ensured the system is production-ready, with **EfficientNet-B0** identified as the optimal classification engine. The final model's outstanding mathbf(97.12%) accuracy demonstrates proficiency in advanced deep learning and its application to satellite imager.

## **9.2. Future Enhancements**

Future work will focus on scaling the system and exploring even more advanced machine learning techniques:

- **Data Integration:** Incorporate additional satellite data sources, such as **Landsat** or **Synthetic Aperture Radar (SAR)** data, to improve classification robustness, especially in challenging conditions like cloud cover.
- **Advanced Architectures:** Experiment with newer, state-of-the-art algorithms such as **Vision Transformer (ViT)** models for image classification to explore potential accuracy improvements over CNNs.
- **Scalable Deployment:** Migrate the Flask API to a scalable cloud platform (e.g., AWS Lambda or Azure Functions) and containerize the application using **Docker** for enhanced robustness and real-world scalability.